

Intro to R for Data Analysis

Kristin Yeager, Manager, Statistical Consulting

February 15, 2021

In this workshop

- What is R, and why should I use it?
- What is RStudio, and why should I use it?
- Preview: The benefits of using R
- Key terms
- Worked example of the R workflow
 - Setting up RStudio project
 - Importing data
 - Data cleaning/data management
 - Data analysis

What is R?

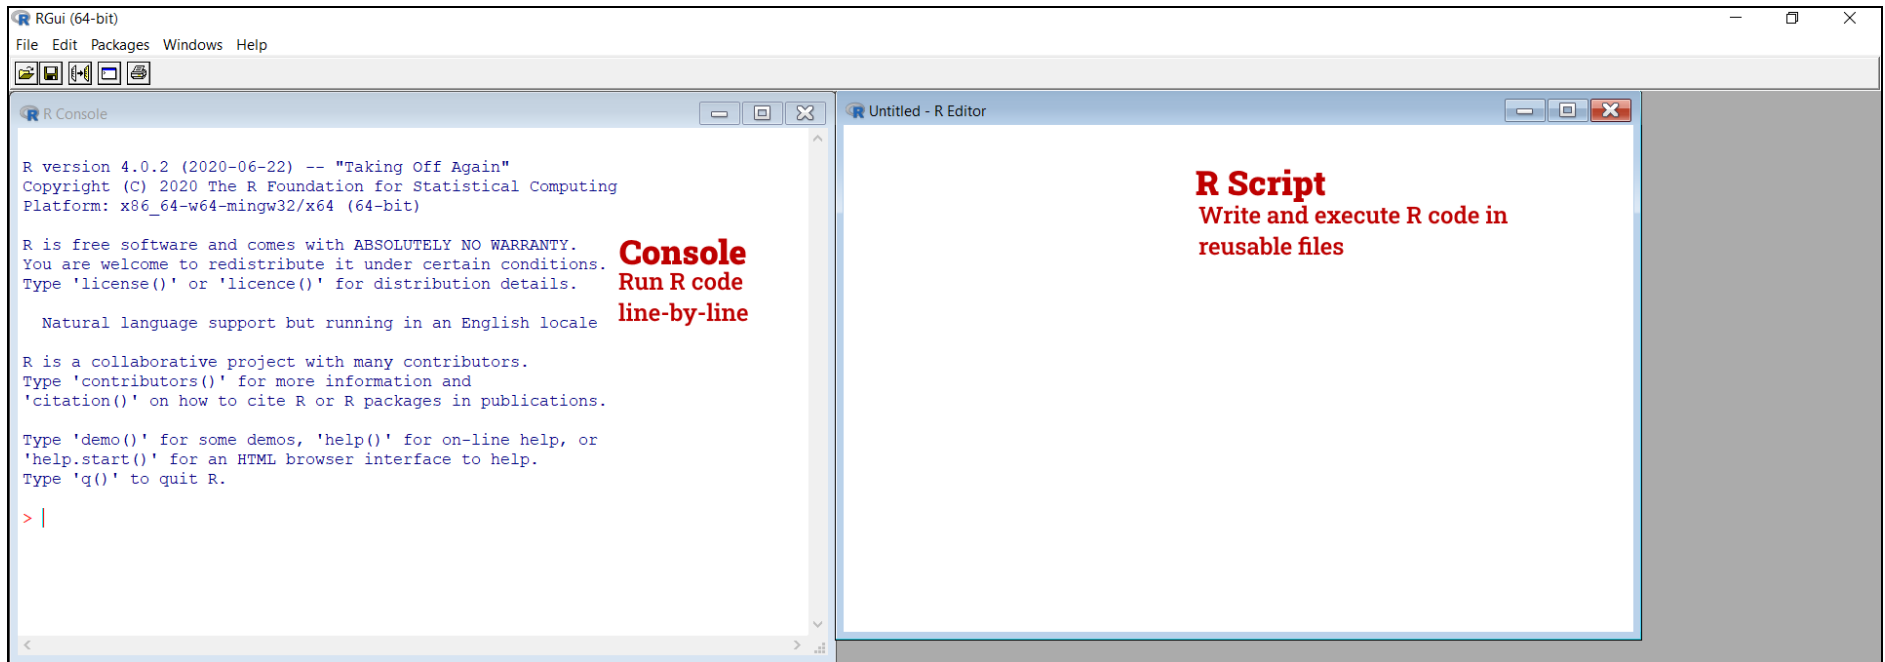
And why you should use it

About R



- What it is: Free and open source statistical software and statistical programming language
- Cost: Free!
- Operating systems: Windows, Mac, Linux
- To install: cran.r-project.org

The R user interface



- The **console** is the “command line” interface part of R – type code into it, press enter, get results back
- The console is fine for “playing around”, but most of the time, you’ll put everything into an **R script** – a file containing executable R code

Why use R?

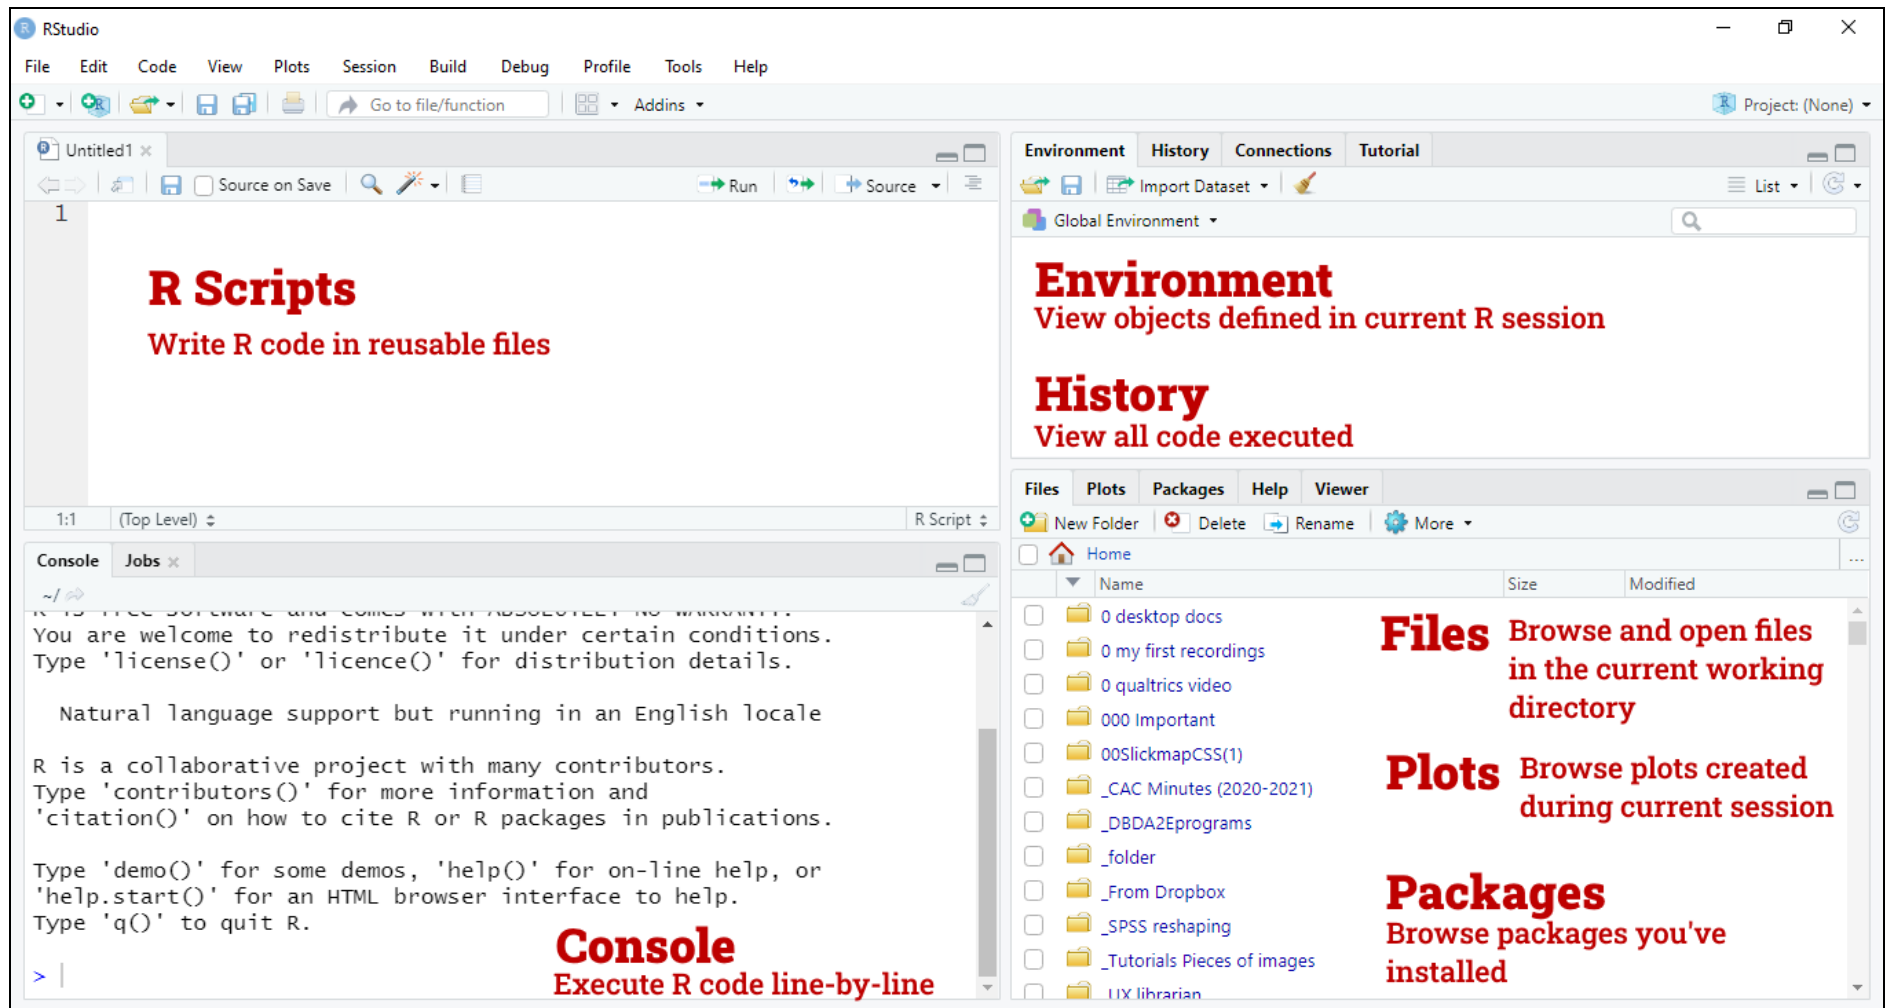
- Free and open-source
- Reproducibility and *literate statistical programming* (combining written report with the code used to generate tables, figures, analysis)
- Beautiful, publication-quality graphics (especially ggplot2)
- Frequently the first place to find new and cutting-edge statistical methods (even before they're available in big-name stat software programs)
- Large and active user community
 - Better free tutorials
 - Active forums, social media community for seeking help
 - Opportunities to promote yourself professionally

About RStudio



- Completely optional but can make writing R code much easier!
- What it is: Interactive development environment (IDE) for R
 - View objects in your R environment
 - Syntax highlighting and auto-completion features
 - Keyboard shortcuts for writing code
 - And much, much more
- Requires R to work – can't be used without first installing R
- Cost: Free
- Operating systems: Windows, Mac, Linux
- To install: rstudio.com/products/rstudio/download

The RStudio User Interface



The R Workflow

1. Obtain data and set up project directory
2. Read/import data into R
3. Data management
 - Check data quality
 - Compute new variables
 - Filter rows
 - Select columns
4. Data analysis
 - Descriptive statistics
 - Plots
 - Statistical tests and models

Preview: The benefits of using R

Fast-forwarding to when you're using R to its full potential

R syntax

Necessary vocabulary before we can get to the “good stuff”

Anatomy of a dataset - Data structure

	A	B	C	D	E	F	G	H	I
1	ids	bday	Gender	Age	Athlete	Smoking	Height	Weight	Mile
2	20055	8/6/1995	0	24	1	1	64.5	127.5	494
3	20075		0		0	1	74	195.5	441
4	20087		0		0	1	68.5	153	
5	20088	1/8/2001	1	18	0		67.5	158.5	514
6	20135	12/26/1998	0	20	0	1	75	152	450
7	20161	12/2/2000	1	18	1	1	73.5	177.5	403
8	20188	8/6/2001	0	18	0	1	76	181	
9	20215	8/5/1997	0	22	0	2	76	155.5	563
10	20250	12/14/1994		24	0	1	66.5	119	653
11	20354	8/5/1999	1	20	0	1	62.5	147.5	684
12	20640	5/31/2000	1	19	0	1	63	161	798
13	20739	6/23/1999	1	20	0	1	72	190.5	612

- First row contains variable names (no spaces or special characters; numbers, periods, underscores OK; R is case-sensitive)
- Rows = observations (1 row = 1 respondent or unit)
- Columns = variables (1 column = 1 distinct variable)

Objects & Environment

- **Objects** are named “things” you’ve created in R – can be manipulated, modified, plotted, etc.
- Objects are created using the assignment operator: `<-`

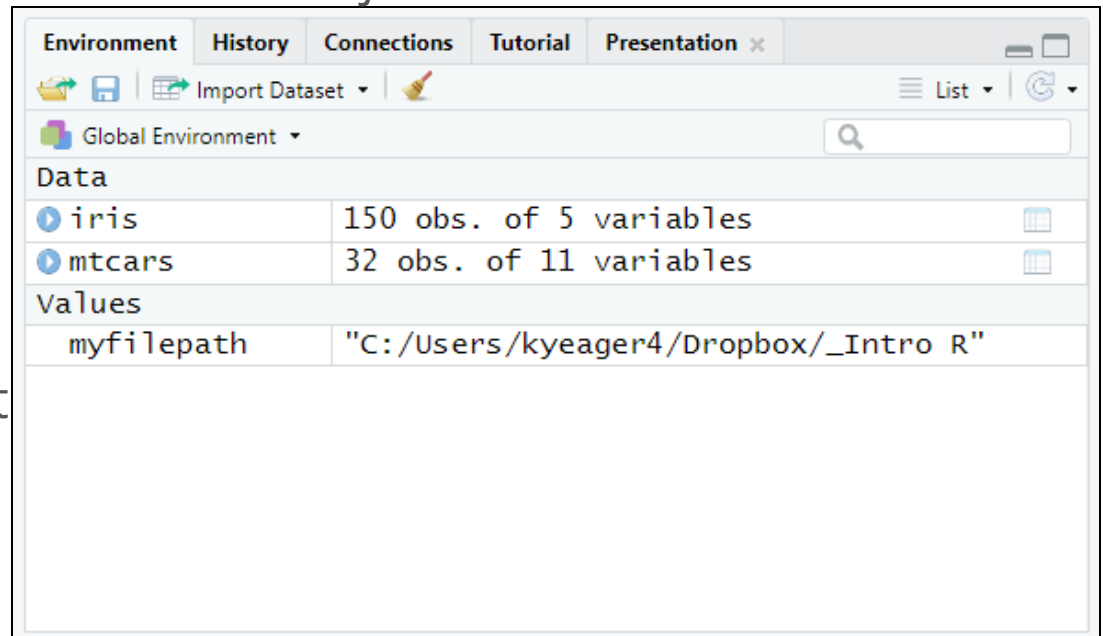
Example: Create a new object called `new_object` whose value is equal to a constant:

```
new_object <- 10
```

There are different types of objects in R – in this session, we’ll be focusing on **dataframes** and **vectors**

Objects & Environment

- The **environment** is the collection of all R objects that have been defined during the current session
- If you're using RStudio, you can view all objects in the Environment panel
- The environment is normally cleared when you close R, but you can choose to save the environment so it still exists the next time you launch R (normally not recommended)



Functions and arguments

- R syntax is generally driven by **functions**, which are pre-programmed tasks that R can do
- Most (but not all) functions have **arguments** that can modify or control the behavior of the function

Example: Compute the mean of a variable

```
mean(x=mydata$SleepTime, na.rm=TRUE)
```

- `mean()` is the name of the function
- `x` and `na.rm` are arguments of `mean()`
 - `x` is the vector of data to compute the mean of
 - `na.rm=TRUE` tells R to ignore any missing values in the input vector

Functions and arguments

- You can open a function's documentation page using the function `help()`:

```
help("mean")  
help("read.csv")
```


R packages

- R packages are user-created modules of code intended to fulfill a narrow task or purpose
- Shared through repositories, of which CRAN is “main” one
 - See [CRAN Task Views](#) to explore R packages by topic
- An R user “installs” an R package from the online repository to their local package directory using the function `install.packages()`
- Once a package has been installed, it can be used in your R scripts – simply load the package using `library()`

Worked Example

Today's sample data

Tutorial sample data:

- Simulated survey dataset we use for our online tutorials
- n=435 “college students”
- Available in CSV, TXT, SAS, SPSS formats
- Download from our website: libguides.library.kent.edu/SPSS/ ([direct download link for CSV](#))

Research questions we'll consider today:

- Do students who study more get less sleep? (Or: do students who sleep more study less?)
- Do underclassmen (freshmen and sophomores) get less sleep than upperclassmen (juniors and seniors)?

Working in R, part 0

Before you open the program

Set up your project

Create a project folder on your computer with the following subdirectories:

- project
 - code
 - data [raw]
 - data ["processed"]
 - plots
 - output
 - reports

RStudio Projects

- Reason to do this: Can declare a folder/directory on your computer as an “RStudio project”
- Allows you to read and write files using relative paths instead of absolute paths

Example: Absolute file path on Windows:

```
"C:/Users/yourusername/yourprojectfolder/data/mydata.csv"
```

Example: Absolute file path on Mac:

```
"/Users/yourusername/yourprojectfolder/data/mydata.csv"
```

Example: Relative file path after creating an RStudio project in the “yourprojectfolder” folder:

```
"data/mydata.csv"
```

Working in R, part 1

Importing text-based data

Step 1: Get data into R

- *Data frames* are R's object type for traditional "data sets"
- Base R function to import text and CSV files as a data frame: `read.table` or `read.csv`

```
mydata <- read.csv(file="data/Sample_Dataset_2014.csv")
```

This code creates an object named `mydata` containing the data from this file as a data frame.

Step 2: Look at the data we imported into R

How do we know if our import was successful?

- View the dataframe as a spreadsheet using `View(mydata)`
- Print variable types using function `str(mydata)` (short for `structure`)
- Print minimum, maximum, median, missing values for all variables in dataframe using `summary(mydata)`
- Print first 5 rows of dataframe using `head(mydata)`
- Print last 5 rows of dataframe using `tail(mydata)`
- Print names of variables in dataframe using `names(mydata)`

Tip

Never use the `attach()` function, even if you see it in an online tutorial and it looks like it would make things easier.

Step 3: Access variables within our dataset

- Variables are a named component of our dataframe
- Some functions expect a dataframe; other functions expect a *vector* of observations – that is, a single, specific column from a dataframe

When we need to access a column of a dataframe as a vector, we use the **\$** operator:

```
mydata$varname
```

This syntax can be combined with the assignment operator to add new variables to a dataframe, or edit existing variables in a dataframe:

```
mydata$varname_times_ten <- mydata$varname*10
```

Step 4: Access rows within our dataset

Normally when we want to access specific rows, what we really want is to filter our data by some condition. We can do this using function `subset()`:

```
freshmen_only <- subset(mydata, Rank==1)
```

The first argument to `subset` is the name of the dataframe. The second argument is a **logical condition** for which rows to keep (here, `Rank==1`, i.e. keep only freshman).

Working in R, part 2

Summary statistics

Summary statistics for continuous numeric variables

- `mean()`, `sd()`, `min()`, `max()`, `median()`, `sum()`
 - The `na.rm=TRUE` argument
- Base R graphics: `hist()`, `boxplot()`

Summary statistics for categorical variables

- `table()`
- `prop.table()`
- Base R graphics: `barplot(table(...))`

Working in R, part 3

Installing an R package for plotting

Installing package ggplot2

- The ggplot2 package creates beautiful, publication-quality graphics using the *grammar of graphics* syntax
- To install the package:

```
install.packages("ggplot2")
```

Loading package ggplot2

After you've installed the package, you can load the package whenever you want to use it:

```
library(ggplot2)
```

We typically put these package loading statements at the start of our scripts.

Working in R, part 4

Data analysis

Examples

- Correlation using `cor`
- Scatterplot using `ggplot2`
- Linear regression using `lm()`

Questions?

Appendix A

R version used for this workshop

R version used for this workshop

```
sessionInfo()
```

```
## R version 4.0.2 (2020-06-22)
## Platform: x86_64-w64-mingw32/x64 (64-bit)
## Running under: Windows 10 x64 (build 18363)
##
## Matrix products: default
##
## locale:
## [1] LC_COLLATE=English_United States.1252
## [2] LC_CTYPE=English_United States.1252
## [3] LC_MONETARY=English_United States.1252
## [4] LC_NUMERIC=C
## [5] LC_TIME=English_United States.1252
##
## attached base packages:
## [1] stats      graphics  grDevices  utils      datasets  methods   base
##
## loaded via a namespace (and not attached):
## [1] compiler_4.0.2  magrittr_2.0.1  tools_4.0.2    htmltools_0.5.0
## [5] yaml_2.2.1      stringi_1.5.3   rmarkdown_2.6  knitr_1.30
## [9] stringr_1.4.0   xfun_0.19       digest_0.6.27  rlang_0.4.10
## [13] evaluate_0.14
```